

MIE1517. Midterm Study Topics

This is a list of lecture/assignment topics compiled by your instructor to help you study for the midterm. Note that the list may not be exhaustive.

In particular, the list of topics doesn't cover the skills that we develop throughout the course through the labs. For example, interpreting the training curve is a skill that we practiced throughout the course. As another example, you should be able to predict the output shape of PyTorch layers we discussed in class, like `nn.Conv2d`.

Machine Learning Overview

- artificial intelligence, machine learning, deep learning
- type of problems where machine learning is successful
- supervised learning (regression, classification), unsupervised learning, reinforcement learning
- issues with deep learning: interpretability, fairness

Biological motivations

- biological neuron: dendrite, axon, synapse, cell body, the direction that information travel

Artificial Neuron

- artificial neuron: activation, weight, bias
- input layer, hidden layer(s), output layer
- fully-connected layer
- feed-forward neural network
- multi-layer perceptron (MLP)
- steps to training a neural network
- loss function (binary cross-entropy)
- optimizer (stochastic gradient descent)
- overfitting, training set, test set, prediction

Artificial Neural Networks

- ReLU, sigmoid, tanh
- neural network architecture
- counting number of layers in an ANN (i.e. MLP)
- parameters, weights, biases
- counting number of parameters / weights / biases
- standard training/test set splits

Optimization and Hyperparameter

- batch size (effect of large / small batch size)
- iteration, epoch
- the use of PyTorch's DataLoader to batch data for us
- learning rate (effect of large / small learning rate)
- training curve (what's on each axis)
- training loss / accuracy / error
- interpreting training curves
- validation set
- validation loss / accuracy / error
- hyperparameter search
- examples of hyperparameters beyond learning rate and batch size
- model selection (based on validation)
- checkpointing

Multi-Class Classification

- the use of PyTorch fully-connected nn.Linear layer
 - the meaning of in_features and out_features
 - changing the number of layers in a neural network
- the use of PyTorch activation functions F.relu and F.softmax
- the difference between PyTorch CrossEntropyLoss and BCEWithLogitsLoss
- the general training loop code:

```
out = model(imgs) # forward pass
loss = criterion(out, labels) # compute the total loss
loss.backward() # backward pass (compute parameter updates)
optimizer.step() # make the updates for each parameter
optimizer.zero_grad() # a clean up step for PyTorch
```

- interpreting training curves to identify overfitting
- forward pass, backwards pass
- predict the output shape of a PyTorch layer

Troubleshooting

- the idea of overfitting on a small data set

Convolutions

- reasons why fully-connected layers can be problematic

- locally-connected layers, weight-sharing
- convolutional kernels / filters and their sizes
- input and output channels / feature maps and their sizes
- convolutional arithmetics (computing forward pass of a convolution layer given the input and kernel)
- PyTorch “NCHW” format
- PyTorch convolution layers `nn.Conv2D`, its parameters, and computing the output shape of a convolution
- zero-padding (the reason for their use, selecting padding size)
- max-pooling, and the `nn.MaxPool2D` layer
- average-pooling
- strided convolutions
- examples of the types of features computable using convolutions
- forward pass arithmetic of pooling and padded/strided convolutions
- feature visualizations and receptive fields
- gradient ascent and style transfer
- image classification, bounding boxes and object detection
- semantic and instance segmentation

Architectures

- What is new in each architectures:
 - LeNet
 - AlexNet
 - GoogLeNet: repeated modules, multiple convolution kernels
 - ResNet: residual connections (i.e. skip connections)

Regularization

- why we don't talk about underfitting as much
- how to detect overfitting
- data normalization
- data augmentation
- weight decay
- early stopping
- dropout
 - marking whether a PyTorch model is in training mode or evaluation mode

Autoencoders

- the idea of an embedding space and different features/representation of the input
- mean-squared error loss (`MSELoss`)

- how to train an autoencoder
- how to use the decoder to generate new images
- the effect of interpolation in the embedding space vs interpolating in the pixel space
- possible uses of an autoencoder
- effect of size of latent space
- semi-supervised learning with autoencoders
- main differences between the different autoencoder architectures

Convolution Transpose

- upsampling
- convolution transpose arithmetic
 - effect of padding
 - effect of strides
 - effect of output_padding